

- ✓ Title : Design and implement Deep Convolutional GAN to generate images of faces/digits from a set of given images.

```
import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
from keras.models import Sequential
from keras.layers import Dense, Flatten, Dropout, BatchNormalization, LeakyReLU, Reshape
from keras.datasets import fashion_mnist
```

- ✓ Load Dataset

```
(train_x, _), (_, _) = fashion_mnist.load_data()
train_x = train_x / 255.0
train_x = train_x.reshape(-1, 28, 28, 1)
train_x = train_x * 2 - 1
print(train_x.shape)
```

```
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-labels-idx1-ubyte.gz
29515/29515 ————— 0s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-images-idx3-ubyte.gz
26421880/26421880 ————— 0s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-labels-idx1-ubyte.gz
5148/5148 ————— 0s 1us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-images-idx3-ubyte.gz
4422102/4422102 ————— 0s 0us/step
(60000, 28, 28, 1)
```

- ✓ Generator

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, LeakyReLU, BatchNormalization, Reshape, Input

generator = Sequential([
    Input(shape=(100,)), # ✓ correct way

    Dense(512),
    LeakyReLU(negative_slope=0.2),
    BatchNormalization(momentum=0.8), # ✓ FIXED

    Dense(256),
    LeakyReLU(negative_slope=0.2),
    BatchNormalization(momentum=0.8),

    Dense(128),
    LeakyReLU(negative_slope=0.2),
    BatchNormalization(momentum=0.8),

    Dense(784, activation='tanh'), # GAN standard
    Reshape((28, 28, 1))
])

generator.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
dense_4 (Dense)	(None, 512)	51,712
leaky_re_lu_3 (LeakyReLU)	(None, 512)	0
batch_normalization_3 (BatchNormalization)	(None, 512)	2,048
dense_5 (Dense)	(None, 256)	131,328
leaky_re_lu_4 (LeakyReLU)	(None, 256)	0
batch_normalization_4 (BatchNormalization)	(None, 256)	1,024
dense_6 (Dense)	(None, 128)	32,896
leaky_re_lu_5 (LeakyReLU)	(None, 128)	0
batch_normalization_5 (BatchNormalization)	(None, 128)	512
dense_7 (Dense)	(None, 784)	101,136
reshape_1 (Reshape)	(None, 28, 28, 1)	0

Total params: 320,656 (1.22 MB)

Discriminator

```
discriminator = Sequential([
    Flatten(input_shape=(28,28,1)),
    Dense(256),
    LeakyReLU(0.2),
    Dropout(0.5),
    Dense(128),
    LeakyReLU(0.2),
    Dropout(0.5),
    Dense(64),
    LeakyReLU(0.2),
    Dropout(0.5),
    Dense(1, activation='sigmoid')
])
discriminator.compile(optimizer='adam', loss='binary_crossentropy')
discriminator.summary()
```

/usr/local/lib/python3.12/dist-packages/keras/src/layers/reshaping/flatten.py:37: UserWarning: Do not pass an `input\_shape` to `input\_shape` in `super().\_\_init\_\_(\*\*kwargs)`

Model: "sequential_2"

Layer (type)	Output Shape	Param #
flatten (Flatten)	(None, 784)	0
dense_8 (Dense)	(None, 256)	200,960
leaky_re_lu_6 (LeakyReLU)	(None, 256)	0
dropout (Dropout)	(None, 256)	0
dense_9 (Dense)	(None, 128)	32,896
leaky_re_lu_7 (LeakyReLU)	(None, 128)	0
dropout_1 (Dropout)	(None, 128)	0
dense_10 (Dense)	(None, 64)	8,256
leaky_re_lu_8 (LeakyReLU)	(None, 64)	0
dropout_2 (Dropout)	(None, 64)	0
dense_11 (Dense)	(None, 1)	65

Total params: 242,177 (946.00 KB)
Trainable params: 242,177 (946.00 KB)

GAN Model

```
discriminator.trainable = False
gan = Sequential([generator, discriminator])
gan.compile(optimizer='adam', loss='binary_crossentropy')
gan.summary()
```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
sequential_1 (Sequential)	(None, 28, 28, 1)	320,656
sequential_2 (Sequential)	(None, 1)	242,177

Total params: 562,833 (2.15 MB)
Trainable params: 318,864 (1.22 MB)

Training

```
epochs=10
batch_size=100
noise_dim=100

for epoch in range(epochs):
    print(f"Epoch {epoch+1}")
    for i in range(train_x.shape[0]//batch_size):
        noise = np.random.normal(0,1,(batch_size,noise_dim))
        fake = generator.predict(noise, verbose=0)
        real = train_x[i*batch_size:(i+1)*batch_size]

        y_real = np.ones((batch_size,1))
        y_fake = np.zeros((batch_size,1))


        discriminator.trainable=True
        discriminator.train_on_batch(real,y_real)
        discriminator.train_on_batch(fake,y_fake)

        noise = np.random.normal(0,1,(batch_size,noise_dim))
        discriminator.trainable=False
        gan.train_on_batch(noise,y_real)

    if epoch%10==0:
        samples=10
        gen = generator.predict(np.random.normal(0,1,(samples,noise_dim)), verbose=0)
        for k in range(samples):
            plt.subplot(2,5,k+1)
            plt.imshow(gen[k].reshape(28,28), cmap='gray')
            plt.xticks([]); plt.yticks([])
        plt.show()

print('Training complete')
```

Epoch 1




Final Output

```
noise = np.random.normal(0,1,(10,100))
gen = generator.predict(noise)

plt.imshow(noise)
plt.title('How the noise looks')
plt.show()

fig, ax = plt.subplots(2,5)
fig.suptitle('Generated Images from Noise using GANs')
idx=0
for i in range(2):
    for j in range(5):
        ax[i,j].imshow(gen[idx].reshape(28,28), cmap='gray')
        ax[i,j].axis('off')
        idx+=1
plt.show()
```

epoch 10
 Training complete 0s 36ms/step



Generated Images from Noise using GANs

